UNITED ARAB EMIRATES
MINISTER OF STATE FOR ARTIFICIAL INTELLIGENCE,
DIGITAL ECONOMY & REMOTE WORK APPLICATIONS OFFICE

Agentic Artificial Intelligence

A National Definitional and Assessment Reference

MAY 2026

This document establishes the official definitions of Generative Artificial Intelligence, Artificial Intelligence Agents, and Multi-Agent Frameworks adopted across UAE government entities, together with the criteria by which entities shall determine whether deployed AI systems qualify as agentic.

Agentic Artificial Intelligence System

An automated software system that manages its operating environment and takes goal driven actions to achieve specific objectives

An Agent: A virtual tool or set of tools that enables users to:

- **Make requests using simple, natural language.**
- **Complete tasks on the user's behalf from the initial request through to achieving the desired outcome.**
- **Fully plan and navigate across different platforms, with the AI agent autonomously accessing multiple portals and payment systems, while remaining under human supervision for monitoring and guidance.**
- **Define a goal and complete it end-to-end by developing the necessary plan and delivering the final result.**

Where today's AI tools tell people what to do, an Agent does it for them within authorised limits, with full audit trails, and under clear human oversight.

Table of Contents

Forward

> Two decades of digital transformation have positioned the United Arab Emirates as a structural transition in the operation of government.

The digital trajectory of the United Arab Emirates has been built across successive milestones: the launch of Electronic Government in 2012, the introduction of Mobile Government (mGov) in 2013, and the UAE Strategy for Artificial Intelligence 2031 in 2017, accompanied by the appointment of the world's first Minister of State for Artificial Intelligence.

These milestones, framed within the UAE Centennial 2071 vision, established the foundation on which the current Programme operates. Successive sector strategies in health, transport, energy, and education have extended that foundation across government.

In April 2026, His Highness Sheikh Mohammed bin Rashid Al Maktoum, Vice President and Prime Minister of the United Arab Emirates and Ruler of Dubai, announced — under the directives of His Highness Sheikh Mohamed bin Zayed Al Nahyan, President of the United Arab Emirates — a new government framework. Within two years, fifty per cent of government sectors, services, and operations will be transformed to operate through agentic Artificial Intelligence. This commitment positions the UAE as the first state to operate at this scale through agentic autonomous systems.

This guide serves as the reference document that demystifies the concept of agentic Artificial Intelligence, unifies its definitions across federal entities, and provides a guide to assist in the classification of their deployed AI systems for the purposes of reporting and evaluation.

Chapter 1
Scope and Purpose

This document establishes a unified national framework for the classification of Artificial Intelligence systems deployed across UAE government entities. It distinguishes among three categories of system: **Generative Artificial Intelligence, Artificial Intelligence Agents**, and **Multi-Agent Frameworks.**

PURPOSE

1. **Standardisation of terminology**. To establish a single, defensible meaning for the term "agentic" for use in policy, procurement, reporting, and operational evaluation across all UAE government entities.
2. **Cut-off criteria for classification**. To provide unambiguous criteria by which a deployed system may be determined to qualify, or to fail to qualify, as agentic for the purposes of the National Programme.
3. **Prevention of misclassification**. To preempt classification errors common in the early phase of agentic AI adoption, including the labelling of conventional chatbots, scripted automations, and presentation interfaces as Agents or Multi-Agent Frameworks.
4. **Support for transformation and strategy**. To provide a mechanism to support government entities in their transformation and strategy-setting processes.

SCOPE OF APPLICATION

This is a definitional and evaluative reference. It is intended for use by UAE government entities, will help in classifying their deployed AI systems, explains the different classes of AI, and provides information for reporting and evaluation. This document is supported by a set of documents addressing governance, procurement, capability development, measurement, and citizen transparency, issued separately by the taskforce. Cybersecurity and operational safety considerations apply to every deployment regardless of category, and are addressed in the supporting documents.

SIGNIFICANCE

Performance under the National Programme will be subject to formal measurement. Clear differentiation between AI agentic systems and the previous generation of AI systems is essential for keeping pace with rapid technological advancements and maintaining global leadership in the field of artificial intelligence.Adopting a traditional chatbot, while assuming that it possesses advanced AI agentic capabilities, can lead to a decline in the quality of the customer experience and a failure to achieve the desired efficiency and user experience Consistent and rigorous classification is therefore a precondition for honest measurement, and honest measurement is a precondition for the calibration of the Programme.

Chapter 2

Generative Artificial Intelligence

REFERENCE DEFINITION

"Generative AI (GenAI) is a category of AI that can create new content such as text, images, videos and music. It gained global attention in 2022 with text-to-image generators and large language models."

— UAE Guidelines on the Use of Generative AI Technologies in Government, ai.gov.ae, 2025. See also: Organisation for Economic Co-operation and Development (OECD), 2024.

A Generative AI system produces content — text, image, audio, or code — in response to a prompt. The system does not act in the operational environment beyond the production of that content, and it does not maintain state across interactions in the absence of an explicit memory architecture. The current generation is built on large language models that produce output by predicting the most statistically probable continuation of a given input.

The capabilities of Generative AI are substantive and form the substrate on which Artificial Intelligence Agents are constructed. The boundary between Generative AI and an Agent is, however, well-defined: a Generative AI system produces content; it does not execute action.

CAPABILITIES

- **Drafting**. Production of first versions of letters, memoranda, briefings, and reports.
- **Summarisation**. Reduction of long-form documents to structured briefings.
- **Translation**. Translation between Arabic, English, and other languages with sensitivity to register and context.
- **Extraction**. Extraction of structured data from unstructured documents.
- **Generation**. Generation of code, scripts, formatted content, and document transformations.
- **Explanation**. Provision of natural-language responses to general-knowledge enquiries where accuracy can be independently verified.

DEFINING CHARACTERISTICS

A system is classified as Generative AI when all of the following conditions are met:

1. **Prompt-driven operation**. The system produces output only in response to a prompt and does not initiate action autonomously.
2. **Single-turn interaction**. Each interaction is materially complete in itself; there is no ongoing task across interactions.
3. **Stateless operation**. The system does not retain context across sessions in the absence of an explicit memory architecture.
4. **Absence of tool invocation**. The system does not invoke external systems, APIs, or data sources as part of its operation.

OPERATIVE LIMIT

A Generative AI system provides information regarding a service; it does not perform the service. At the point at which the system is granted the authority and the tools required to act on a defined goal, it constitutes a different category of system, addressed in Chapter Three.

Chapter 3
Artificial Intelligence Agents

On the term "Agentic"

Agentic describes any system that satisfies the four-property test below. An individual Agent is agentic in itself; a Multi-Agent Framework (Chapter Four) is a coordinated arrangement of agentic Agents under an orchestration layer.

REFERENCE DEFINITION

> "An artificial intelligence (AI) agent is a system or program that is capable of autonomously performing tasks on behalf of a user or another system by designing its workflow and utilizing available tools."
>
> — IBM, What are AI agents?, 2024

An Artificial Intelligence Agent is a Generative AI system extended with three additional capabilities: a defined goal, the authority to act, and access to a set of tools through which action may be taken. The Agent operates by repeating a simple sequence: it considers what to do next, takes an action through one of its tools, observes the result, and advances to the next step. This sequence continues until the goal has been achieved, or until the Agent determines it cannot proceed. Where a Generative AI system responds to a single prompt, an Agent receives an objective — for example, submitting an application on a citizen's behalf by retrieving the relevant records, completing the form, and uploading the required documents — and proceeds through the work without continuous re-prompting at each step.

CAPABILITIES

- **Decomposition**. Receipt of a defined outcome and decomposition of that outcome into the intermediate steps required for its achievement.
- **Action execution**. Invocation of external systems, APIs, and data sources, including for the purposes of writing, submitting, paying, scheduling, and notifying.
- **Iteration**. Repetition of the consider–act–observe sequence until the goal is achieved or until the Agent determines it cannot proceed within its authority.
- **State maintenance**. Retention of task state across multiple steps and, where designed for it, across sessions.

THE FOUR PROPERTIES OF AN AGENT

A system is classified as an Agent when all four of the following conditions are met. The absence of any one property results in classification under a different category.

1. **Goal-directedness**. The system is provided with an outcome rather than a script.
2. **Tool use**. The system invokes functions external to itself for the execution of action in the operational environment.
3. **State maintenance**. The system maintains context across the steps of a task, and where designed for it, across sessions.
4. **Iterative execution**. The Agent autonomously repeats the consider–act–observe sequence until the goal is achieved or it determines it cannot proceed.

OPERATIVE SCOPE

An Artificial Intelligence Agent does not provide information regarding a service; it executes the service, in whole or in part, under an authorised scope and oversight regime.

Chapter 4
Multi-Agent Frameworks

INDUSTRY FRAMING · IBM

IBM defines AI agent orchestration as "the process of coordinating multiple specialized AI agents within a unified system." A central orchestrator — which may itself be an agent or a framework — directs the work of the specialised agents toward a shared goal.

— IBM, What is AI Agent Orchestration?, 2026

INDUSTRY FRAMING · MCKINSEY

McKinsey describes the orchestration layer as "a composable, distributed, and vendor-agnostic architectural paradigm." In this paradigm, agents work as a coordinated network across the organisation's diverse systems, tools, and language models.

— McKinsey & Company, Seizing the agentic AI advantage, 2025

THE ORCHESTRATION LAYER

1. **Planning**. Translation of a high-level outcome — for example, the renewal of a driving licence — into the sequence of agent-level tasks required to achieve it.
2. **Delegation**. Routing of each sub-task to the Agent best suited to its execution by domain, by tool authorisation, or by authority level.
3. **Supervision**. Monitoring of Agent performance, enforcement of budget, time, and step limits, and termination of operation upon guardrail breach.
4. **Consolidation**. Receipt of outputs from multiple Agents and assembly of those outputs into a single coherent result for the citizen, the operator, or the calling system.

MULTI-AGENT SYSTEMS

A Multi-Agent Framework deployed in government is typically a coordinated set of specialised Agents — each with a domain-specific context and a more constrained tool inventory than a generalist Agent — operating under a single orchestration layer.

CRITERIA FOR CLASSIFICATION AS A MULTI-AGENT FRAMEWORK

1. At least one component meets the four-property test for classification as an Agent (Chapter Three).
2. An orchestration layer is present, with explicit responsibility for planning, delegation, supervision, and consolidation.
3. The framework is constrained by **deterministic guardrails** that operate independently of the Agents themselves.
4. Every Agent action is **logged, auditable, and reversible** at the level required by the use case.
5. A **named individual** is accountable for the framework as a whole, in addition to any accountability assigned for individual Agents.

COMMON MISCLASSIFICATION

A deployment of multiple Agents without coordination between them is not a Multi-Agent Framework — it is a collection of independent Agents, each agentic in its own right. Coordination through an orchestration layer is what makes the collection a Framework, and what makes the Framework governable.

Chapter 5

Comparative Analysis

The following table sets out the principal points of distinction among the three categories across nine operational dimensions.

DIMENSION	Generative AI	AI Agent	Multi-Agent Framework
	(LLM-based assistant)	(single goal, single workflow)	(orchestration of multiple Agents)
INITIATION	Awaits a prompt	Receives a goal	Receives an outcome
STEPS	Single response	Multi-step loop within a task	Multi-Agent, multi-step workflow
TOOL INVOCATION	None	Bounded set; one Agent	Distributed across Agents
STATE RETENTION	Stateless by default	Across steps; sometimes sessions	Persistent and cross-Agent
OUTPUT	Content	Content or action	End-to-end outcome
AUTHORITY	None; advisory only	Bounded by defined scope	Bounded by both Agents and orchestrator
COORDINATION	Not applicable	Single Agent	Orchestrator and specialised Agents
SUITED TO	Drafting, summarisation, enquiry	Defined process or sub-process	Cross-functional service journeys
OVERSIGHT	Light review of output	Logged; gated at high-stakes points	Orchestrator, per-Agent, system-wide

Non-Agentic AI

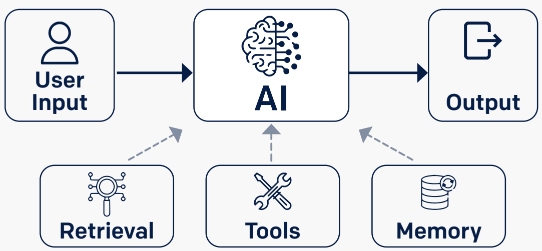

The model can access information, tools and memory, but it mainly responds to prompts rather than inde-pendently purcuing a goal.

AI Agent

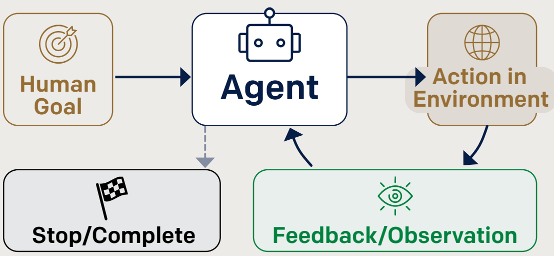

The system chooses actions, observes results, updates its approach, and continues until the goal is achieved or it cannot proceed.

OBSERVATION

Operational capability and operational risk increase together. Each progression from Generative AI to Agent to Multi-Agent Framework adds both functional capability and exposure. Governance does not arise automatically from the deployment of a more capable system; it is a cost the deploying entity must pay. A Multi-Agent Framework that lacks proportionate governance presents a greater operational risk than no framework at all.

Chapter 6

Illustrative Application: Driving Licence Renewal

The following section presents the same government service — the renewal of a driving licence — as it would be delivered through each of the three categories of system.

A. GENERATIVE AI · INFORMATION PROVISION

"How do I renew my driving licence?"

The system returns a clear, structured explanation of the renewal procedure: the relevant authority, the documents required, the cost, the means of payment, and the locations at which renewal may be effected. The execution of the service remains with the citizen logging in to the relevant portal, completing the form, uploading the photograph, settling fines, scheduling the appointment, and confirming the renewal.

- Information delivered.
- No identity verification.
- No record retrieval.
- No payment processed.
- No renewal performed.

The Generative AI system informs the citizen; it does not execute the service.

B. AI AGENT · PARTIAL EXECUTION

"Help me renew my driving license."

The Agent verifies the citizen's identity through UAE Pass, retrieves the licence record, confirms eligibility for renewal, settles outstanding traffic fines from the linked payment instrument, and schedules an eye-test appointment at the citizen's nearest centre.

The Agent then requests the inputs that require direct citizen action — for example, a current photograph and authorisation of the renewal payment — and provides instructions for their submission.

- Identity verified.
- Record retrieved.
- Fines settled.
- Appointment scheduled.
- Citizen steps remain.

A small number of citizen actions remain; the structured work of the service has been executed.

C. MULTI-AGENT FRAMEWORK · END-TO-END EXECUTION

"Renew my licence."

The orchestration layer delegates the request to a set of specialised Agents and supervises their joint execution.

- **Identity Agent** verifies the citizen via UAE Pass.
- **Records Agent** retrieves the licence record and confirms eligibility and dues.
- **Fines Agent** presents and settles outstanding fines.
- **Appointments Agent** books the eye-test at the nearest available slot in the calendar.
- **Communications Agent** conducts the interactions in the citizen's preferred language and channel.

A request, a single coherent response, and a renewed licence. No tabs opened, no forms completed, no re-entry of known information; the service is performed end-to-end.

SUMMARY OF PROGRESSION

In the progression from Generative AI to Agent to Multi-Agent Framework, the locus of effort shifts from the citizen to the system. Four distinct interactions across three entities are reduced to a single coordinated process. The Generative AI system informs the citizen. The Agent executes part of the service. The Multi-Agent Framework executes the service end-to-end. This shift, measured at the level of services and workflows across federal entities, constitutes the substance of the National Programme.

Chapter 7

Self-Assessment

A Five-question assessment for the classification of an AI deployment within a UAE government entity.

A system is agentic if it answers "Yes" to all five.

1	**Goal persistence**	Once given an objective, does the system continue working toward it across multiple steps without requiring a new instruction for each step?	☐ Yes ☐ No
2	**Autonomous decomposition**	Does the system decide on its own what steps to take, and in what order, to achieve the objective?	☐ Yes ☐ No
3	**Action capability**	Can the system take actions that affect state outside itself, invoking tools, APIs, or systems that do more than return information to the conversation?	☐ Yes ☐ No
4	**Closed-loop iteration**	Does the system observe the result of each action and use that observation to decide its next action, continuing until the objective is met or it determines it cannot proceed?	☐ Yes ☐ No
5	**Working memory**	Does the system carry context forward across the steps of a task, so that later steps can build on what earlier steps produced?	☐ Yes ☐ No

If any answer is "No," the system may be AI-enabled, but it is not fully agentic by this test.

Chapter 8

Frequently Asked Questions

The following questions address the most common classification queries raised by entities evaluating their existing AI deployments against the criteria established in this reference.

The entity has deployed a chatbot. Does it qualify as agentic?

NOT BY DEFAULT. If it mainly answers questions or retrieves information, it is still a Generative AI application. It becomes an Agent only when it can pursue a goal, use tools, and complete bounded work on the user's behalf across multiple steps — to complete a service or a process. It should only be counted as a Multi-Agent Framework when those actions operate inside a governed orchestration layer with logs, controls, and ownership.

The entity uses Robotic Process Automation (RPA) to automate routine workflows. Does this qualify?

NO. RPA follows a fixed sequence of steps that a human has scripted in advance. An Agent determines its next action within an authorised procedure based on what it observes, not from a pre-defined script. The two can work together — an Agent can call RPA scripts as one of its tools — but a system that simply repeats a pre-written sequence is automation, not agency.

A smart avatar conducts citizen interactions in Arabic and English. Does this qualify as agentic?

NOT IN ITSELF. The avatar itself is only the interface. A voice, face, or embodied experience does not change the classification. What matters is whether the underlying system evaluates, plans, uses tools, acts across steps, and is governed appropriately.

The entity has deployed a "co-pilot" that drafts replies for staff. Does this qualify?

NO. A co-pilot that drafts replies for staff is Generative AI — the staff member remains the one acting. The deployment becomes agentic only when the system can send the reply, update the record, and take the next step on its own, within an authorised scope and oversight regime.

An Agent in the entity is connected to a single tool — it can send an email. Is one tool sufficient?

CONDITIONAL. Having a tool is necessary, but not sufficient. A system that reads one message and sends one reply is Generative AI with a delivery channel — not an Agent. An Agent determines, within its authorised procedure, whether to send, what to send, when to send, and what to do next based on the response. One tool can be enough; a single such determination is not.

Multiple Agents operate within the entity, each handling its own task. Does this constitute a Multi-Agent Framework?

NOT YET. A Multi-Agent Framework needs an orchestration layer that coordinates the Agents — planning the work, delegating sub-tasks, supervising execution, and consolidating results. Several Agents working in parallel, however well-built individually, do not make a Framework. The orchestration layer is what turns a collection of Agents into a coordinated system that delivers an end-to-end service.

Does multi-agent always mean better?

NO. Multi-agent designs are justified when the task genuinely requires specialist roles or cross-domain coordination. Otherwise they add cost, latency, and failure complexity, and a simpler design is often the better public-sector choice.

Should all processes be agentified?

NO. Agentic AI is best suited to multi-step services where the path is variable but the outcome is bounded — where an Agent can pursue a goal, use tools, and complete work on the citizen's behalf. Other processes are better served by Generative AI, by automation, by simple workflows, or by direct human action. The right choice is driven by the nature of the task, not by the appeal of the technology.

END OF REFERENCE · EDITION 1.0

Definitional Summary

Generative Artificial Intelligence produces content. **An Artificial Intelligence Agent** executes a defined goal through tool use, iteration, and the maintenance of state. **A Multi-Agent** Framework coordinates multiple Agents through an orchestration layer to deliver an outcome end-to-end. A deployment that does not meet these definitions does not qualify as agentic for the purposes of the National Programme on Agentic AI.

ISSUED BY

UNITED ARAB EMIRATES
MINISTER OF STATE FOR ARTIFICIAL INTELLIGENCE,
DIGITAL ECONOMY & REMOTE WORK APPLICATIONS OFFICE

الإمارات العربية المتحدة
مكتب وزير دولة للذكاء الاصطناعي
والاقتصاد الرقمي وتطبيقات العمل عن بعد

COMPANION DOCUMENTS

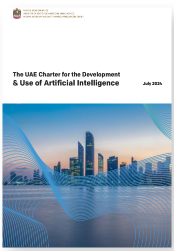

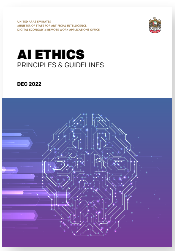

The UAE Charter for the Development & Use of Artificial Intelligence

AI Ethics Guide

Guide to the Use of Generative AI in the UAE Government.

Additional documentation available at ai.gov.ae/publications